

TEJASWINI S

Boston, MA | tejaswinimasters28@gmail.com | [LinkedIn](#) | [Kaggle](#) | [GitHub](#)

SUMMARY

Data scientist with 6 years of experience across data science, machine learning, and QA — most recent 4 years building production ML at Kantar and 2 years earlier as a QA engineer at Wipro. I've worked on Marketing Mix Modelling across many FMCG categories and global markets, helping teams forecast sales and plan pricing, distribution, and media spend. Currently doing my master's in computer science at Northeastern University, and I'm looking for co-op and full-time roles in data science and Machine Learning.

SKILLS & EXPERTISE

Languages: Python, SQL, Core Java

Machine Learning: Regression, Classification, Ensemble Methods (Random Forest, Gradient Boosting, XGBoost, AdaBoost), Feature Engineering, Model Evaluation, Time-Series Forecasting

AI: Supervised & Unsupervised Learning, Neural Networks, NLP, Language Models (LM/LLM)

Libraries & Tools: scikit-learn, PyTorch, pandas, NumPy, seaborn, matplotlib

Cloud: AWS (S3, EC2, Lambda, RDS, DynamoDB)

Databases: MySQL

Visualization & Analytics: Power BI, Marketing Mix Modelling (MMM), Consumer Analytics

WORK EXPERIENCE

Kantar — Senior Analyst, Data Strategy & Engineering Bengaluru, India | Feb 2022 – Dec 2025

- Built and shipped production ML pipelines in Python for Marketing Mix Modelling across 50+ FMCG categories and 15+ global markets, generating long year forecasts that informed pricing, distribution, and media-investment decisions for Fortune 500 clients.
- Engineered multivariate regression and ensemble models (Random Forest, Decision Trees, Gradient Boosting) on 10+ feature pipelines integrating pricing, distribution, media spend, and macroeconomic signals; validated through statistical diagnostics and elasticity constraints to ensure production reliability.
- Developed a client-facing simulator tool serving real-time forecasts and scenario predictions to marketing and strategy stakeholders, transforming offline models into an interactive decision-support product.
- Designed reusable Python automation frameworks and templates that reduced model delivery time and processing errors by 40%, standardizing the modelling workflow across the global forecasting team.
- Recognized with the Rise Team Spark Award 2024 for engineering contributions to Unilever's global Homecare forecasting program; led pilot ML delivery for Kantar's Middle East market entry (GoodyCo, Saudi Arabia).

Wipro — Project Engineer, QA Testing Bengaluru, India | Mar 2017 – Mar 2019

- Performed manual software testing across the full testing lifecycle — test planning, design, execution, and defect tracking — translating business requirements into structured test scenarios and cases.
- Led a team of two testers conducting unit testing across mobile and web interfaces, ensuring quality and functionality of client-facing platforms. Tools: Salesforce, Manual Testing, Core Java.

PROJECTS

Cognate & False-Friend Detection for Dravidian Languages (Kannada–Telugu) 2026

- Developed an NLP classification pipeline to label Kannada–Telugu word pairs as cognate / false friend / unrelated, combining orthographic, phonetic, and semantic features — with multilingual sentence embeddings (LaBSE/MuRIL) as the key signal separating true cognates from false friends.
- Built a 3-way classifier (logistic regression / SVM / MLP) over a concatenated feature vector — edit distance & LCS (orthographic), ALINE/LingPy sound-alignment (phonetic), and embedding cosine similarity (semantic) — benchmarked against an edit-distance baseline.

- Built the candidate dataset by mining ~226K word pairs from IndoWordNet (pyiwn), applying ISO-15919 transliteration and text normalization, with a dual-stream strategy to surface the rare false-friend class.
- Evaluated models with macro-F1, confusion matrices, and feature ablations to isolate each signal's contribution and handle class imbalance.

Global Country-wise Electricity Demand Forecasting

2022

- Forecast monthly national electricity demand (GWh) across 52 countries by comparing six ML models (Linear Regression, Decision Tree, Random Forest, AdaBoost, XGBoost, and a neural network); XGBoost achieved the lowest test RMSE.
- Built the pipeline end-to-end: cleaned ~120K IEA records to 43,500, engineered temporal and country features with one-hot encoding and scaling, and evaluated via RMSE on a held-out test set.

Future Sales Prediction ([Link](#)) — Built regression models (Linear Regression, Decision Tree, Gradient Boosting, Random Forest) to forecast next-month sales per product and store; achieved RMSE of 1.2987 with Gradient Boosting.

Stroke Prediction ([Link](#)) — Classification models to predict stroke risk from patient health data.

Countries for HELP International ([Link](#)) — Unsupervised clustering to categorize countries by development need for aid allocation.

EDUCATION

Northeastern University, Boston, MA

Jan 2026 – Present

Master of Science in Computer Science

Jawaharlal Nehru Technological University, Anantapur, India

2016

Bachelor of Technology in Electronics & Communication Engineering

CERTIFICATIONS

Post Graduate Diploma in Data Science — Manipal Academy of Higher Education, Bangalore, India

Feb 2022

Advanced Certificate in Data Science — IIIT Bangalore, India

Dec 2024